

# Credit Risk-Random forest, NaiveBayes, LDA

April 29, 2019

```
In [28]: # numpy and pandas for data manipulation
import numpy as np
import pandas as pd

# sklearn preprocessing for dealing with categorical variables
from sklearn.preprocessing import LabelEncoder

# File system manangement
import os

# Suppress warnings
import warnings
warnings.filterwarnings('ignore')

# matplotlib and seaborn for plotting
import matplotlib.pyplot as plt
import seaborn as sns

In [29]: # Training data
app_train = pd.read_csv(r'C:\Users\roshn\Desktop\all\application_train.csv')
print('Training data shape: ', app_train.shape)
app_train.head()
```

Training data shape: (307511, 225)

```
Out[29]:
```

	SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE.Cash_loans	\
0	100002.0	1	1	
1	100003.0	0	1	
2	100004.0	0	0	
3	100006.0	0	1	
4	100007.0	0	1	

  

	NAME_CONTRACT_TYPE.Revolving_loans	CODE_GENDER.F	CODE_GENDER.M	\
0	0	0	1	
1	0	1	0	
2	1	0	1	
3	0	1	0	

4			0	0	1
---	--	--	---	---	---

  

	FLAG_OWN_CAR.N	FLAG_OWN_CAR.Y	CNT_CHILDREN	AMT_CREDIT \
0	1	0	0	406597.5
1	1	0	0	1293502.0
2	0	1	0	135000.0
3	1	0	0	312682.5
4	1	0	0	513000.0

  

	...	PA_CNT_DAYS_FIRST_DUE \
0	...	-565.0
1	...	-3823.0
2	...	-784.0
3	...	364266.0
4	...	-6316.0

  

	PA_CNT_DAYS_LAST_DUE_1ST_VERSION	PA_AVG_DAYS_LAST_DUE_1ST_VERSION \
0	125.0	125.000000
1	-3013.0	-1004.333333
2	-694.0	-694.000000
3	366336.0	91584.000000
4	-4186.0	-837.200000

  

	PA_AVG_DAYS_LAST_DUE	PA_AVG_DAYS_TERMINATION \
0	-25.000000	-17.000000
1	-1054.333333	-1047.333333
2	-724.000000	-714.000000
3	182477.500000	182481.750000
4	72136.200000	72143.800000

  

	PA_CNT_HOUR_APPR_PROCESS_START	PA_AVG_HOUR_APPR_PROCESS_START \
0	1.0	9.000000
1	3.0	14.666667
2	1.0	5.000000
3	9.0	14.666667
4	6.0	12.333333

  

	PA_CNT_CNT_PAYMENT	PA_AVG_CNT_PAYMENT	PA_CNT_NFLAG_INSURED_ON_APPROVAL
0	24.0	24.000000	0.0
1	30.0	10.000000	2.0
2	4.0	4.000000	0.0
3	138.0	23.000000	0.0
4	124.0	20.666667	3.0

[5 rows x 225 columns]

In [30]: # Testing data features

app\_test = pd.read\_csv(r'C:\Users\roshn\Desktop\all\application\_test.csv')

```
print('Testing data shape: ', app_test.shape)
app_test.head()
```

Testing data shape: (48744, 224)

```
Out[30]: SK_ID_CURR  NAME_CONTRACT_TYPE.Cash_loans  \
0      100001.0      1
1      100005.0      1
2      100013.0      1
3      100028.0      1
4      100038.0      1

NAME_CONTRACT_TYPE.Revolving_loans  CODE_GENDER.F  CODE_GENDER.M  \
0      0      1      0
1      0      0      1
2      0      0      1
3      0      1      0
4      0      0      1

FLAG_OWN_CAR.N  FLAG_OWN_CAR.Y  CNT_CHILDREN  AMT_CREDIT  AMT_GOODS_PRICE  \
0      1      0      0      568800.0      450000.0
1      1      0      0      222768.0      180000.0
2      0      1      0      663264.0      630000.0
3      1      0      2      1575000.0      1575000.0
4      0      1      1      625500.0      625500.0

...      PA_CNT_DAYS_FIRST_DUE  \
0      ...      -1709.0
1      ...      -706.0
2      ...      -3017.0
3      ...      -3813.0
4      ...      -787.0

PA_CNT_DAYS_LAST_DUE_1ST_VERSION  PA_AVG_DAYS_LAST_DUE_1ST_VERSION  \
0      -1499.0      -1499.000000
1      -376.0      -376.000000
2      -1547.0      -515.666667
3      363664.0      121221.333300
4      -457.0      -457.000000

PA_AVG_DAYS_LAST_DUE  PA_AVG_DAYS_TERMINATION  \
0      -1619.000000      -1612.000000
1      -466.000000      -460.000000
2      -715.666667      -710.333333
3      121171.333300      121182.666700
4      -457.000000      -449.000000
```

	PA_CNT_HOUR_APPR_PROCESS_START	PA_AVG_HOUR_APPR_PROCESS_START	\
0		1.0	13.0
1		2.0	10.5
2		4.0	14.5
3		5.0	10.8
4		2.0	5.5

  

	PA_CNT_CNT_PAYMENT	PA_AVG_CNT_PAYMENT	PA_CNT_NFLAG_INSURED_ON_APPROVAL
0	8.0	8.000000	0.0
1	12.0	12.000000	0.0
2	52.0	17.333333	1.0
3	34.0	11.333333	0.0
4	48.0	24.000000	0.0

[5 rows x 224 columns]

In [31]: *# Create a label encoder object*

```
le = LabelEncoder()
```

```
le_count = 0
```

```
# Iterate through the columns
```

```
for col in app_train:
```

```
    if app_train[col].dtype == 'object':
```

```
        # If 2 or fewer unique categories
```

```
        if len(list(app_train[col].unique())) <= 2:
```

```
            # Train on the training data
```

```
            le.fit(app_train[col])
```

```
            # Transform both training and testing data
```

```
            app_train[col] = le.transform(app_train[col])
```

```
            app_test[col] = le.transform(app_test[col])
```

```
        # Keep track of how many columns were label encoded
```

```
        le_count += 1
```

```
print('%d columns were label encoded.' % le_count)
```

0 columns were label encoded.

In [32]: *# one-hot encoding of categorical variables*

```
app_train = pd.get_dummies(app_train)
```

```
app_test = pd.get_dummies(app_test)
```

```
print('Training Features shape: ', app_train.shape)
```

```
print('Testing Features shape: ', app_test.shape)
```

-----

MemoryError

Traceback (most recent call last)

```
<ipython-input-32-4c395bbbf95> in <module>()
    1 # one-hot encoding of categorical variables
----> 2 app_train = pd.get_dummies(app_train)
    3 app_test = pd.get_dummies(app_test)
    4
    5 print('Training Features shape: ', app_train.shape)

~\Anaconda3\lib\site-packages\pandas\core\reshape\reshape.py in get_dummies(data, prefix, prefix_sep, dummy_na, sparse, drop_first, dtype)
    878         dummy = _get_dummies_1d(col[1], prefix=pre, prefix_sep=sep,
    879                                 dummy_na=dummy_na, sparse=sparse,
--> 880                                 drop_first=drop_first, dtype=dtype)
    881         with_dummies.append(dummy)
    882         result = concat(with_dummies, axis=1)

~\Anaconda3\lib\site-packages\pandas\core\reshape\reshape.py in _get_dummies_1d(data, prefix, prefix_sep, dummy_na, sparse, drop_first, dtype)
    966
    967     else:
--> 968         dummy_mat = np.eye(number_of_cols, dtype=dtype).take(codes, axis=0)
    969
    970         if not dummy_na:

~\Anaconda3\lib\site-packages\numpy\lib\twodim_base.py in eye(N, M, k, dtype, order)
    184     if M is None:
    185         M = N
--> 186     m = zeros((N, M), dtype=dtype, order=order)
    187     if k >= M:
    188         return m
```

MemoryError:

```
In [26]: train_labels = app_train['TARGET']
```

```
# Align the training and testing data, keep only columns present in both dataframes
app_train, app_test = app_train.align(app_test, join = 'inner', axis = 1)
```

```
# Add the target back in
app_train['TARGET'] = train_labels
```

```
print('Training Features shape: ', app_train.shape)
print('Testing Features shape: ', app_test.shape)# Create an anomalous flag column
app_train['DAYS_EMPLOYED_ANOM'] = app_train["DAYS_EMPLOYED"] == 365243
```

```
# Replace the anomalous values with nan
app_train['DAYS_EMPLOYED'].replace({365243: np.nan}, inplace = True)

app_train['DAYS_EMPLOYED'].plot.hist(title = 'Days Employment Histogram');
plt.xlabel('Days Employment');
```

---

```
KeyError                                Traceback (most recent call last)
```

```
~\Anaconda3\lib\site-packages\pandas\core\indexes\base.py in get_loc(self, key, method)
3062         try:
-> 3063             return self._engine.get_loc(key)
3064         except KeyError:
```

```
pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.get_loc()
```

```
pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.get_loc()
```

```
pandas\_libs\hashtable_class_helper.pxi in pandas._libs.hashtable.PyObjectHashTable.get
```

```
pandas\_libs\hashtable_class_helper.pxi in pandas._libs.hashtable.PyObjectHashTable.get
```

```
KeyError: 'TARGET'
```

During handling of the above exception, another exception occurred:

```
KeyError                                Traceback (most recent call last)
```

```
<ipython-input-26-b8da99b4956e> in <module>()
----> 1 train_labels = app_train['TARGET']
      2
      3 # Align the training and testing data, keep only columns present in both dataframes
      4 app_train, app_test = app_train.align(app_test, join = 'inner', axis = 1)
      5
```

```
~\Anaconda3\lib\site-packages\pandas\core\frame.py in __getitem__(self, key)
2683         return self._getitem_multilevel(key)
```

```

2684         else:
-> 2685             return self._getitem_column(key)
2686
2687     def _getitem_column(self, key):

~\Anaconda3\lib\site-packages\pandas\core\frame.py in _getitem_column(self, key)
2690         # get column
2691         if self.columns.is_unique:
-> 2692             return self._get_item_cache(key)
2693
2694         # duplicate columns & possible reduce dimensionality

~\Anaconda3\lib\site-packages\pandas\core\generic.py in _get_item_cache(self, item)
2484         res = cache.get(item)
2485         if res is None:
-> 2486             values = self._data.get(item)
2487             res = self._box_item_values(item, values)
2488             cache[item] = res

~\Anaconda3\lib\site-packages\pandas\core\internals.py in get(self, item, fastpath)
4113
4114         if not isna(item):
-> 4115             loc = self.items.get_loc(item)
4116         else:
4117             indexer = np.arange(len(self.items))[isna(self.items)]

~\Anaconda3\lib\site-packages\pandas\core\indexes\base.py in get_loc(self, key, method)
3063         return self._engine.get_loc(key)
3064     except KeyError:
-> 3065         return self._engine.get_loc(self._maybe_cast_indexer(key))
3066
3067     indexer = self.get_indexer([key], method=method, tolerance=tolerance)

pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.get_loc()

pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.get_loc()

pandas\_libs\hashtable_class_helper.pxi in pandas._libs.hashtable.PyObjectHashTable.get

pandas\_libs\hashtable_class_helper.pxi in pandas._libs.hashtable.PyObjectHashTable.get

```

KeyError: 'TARGET'

```
In [7]: anom = app_train[app_train['DAYS_EMPLOYED'] == 365243]
        non_anom = app_train[app_train['DAYS_EMPLOYED'] != 365243]
        print('The non-anomalies default on %0.2f%% of loans' % (100 * non_anom['TARGET'].mean))
        print('The anomalies default on %0.2f%% of loans' % (100 * anom['TARGET'].mean()))
        print('There are %d anomalous days of employment' % len(anom))
```

The non-anomalies default on 8.07% of loans

The anomalies default on nan% of loans

There are 0 anomalous days of employment

```
In [8]: app_test['DAYS_EMPLOYED_ANOM'] = app_test["DAYS_EMPLOYED"] == 365243
        app_test["DAYS_EMPLOYED"].replace({365243: np.nan}, inplace = True)

        print('There are %d anomalies in the test data out of %d entries' % (app_test["DAYS_EMPLOYED_ANOM"].sum(), app_test["DAYS_EMPLOYED"].count))
```

There are 9274 anomalies in the test data out of 48744 entries

```
In [9]: # Find correlations with the target and sort
        correlations = app_train.corr()['TARGET'].sort_values()

        # Display correlations
        print('Most Positive Correlations:\n', correlations.tail(15))
        print('\nMost Negative Correlations:\n', correlations.head(15))
```

Most Positive Correlations:

OCCUPATION_TYPE_Laborers	0.043019
FLAG_DOCUMENT_3	0.044346
REG_CITY_NOT_LIVE_CITY	0.044395
FLAG_EMP_PHONE	0.045982
NAME_EDUCATION_TYPE_Secondary / secondary special	0.049824
REG_CITY_NOT_WORK_CITY	0.050994
DAYS_ID_PUBLISH	0.051457
CODE_GENDER_M	0.054713
DAYS_LAST_PHONE_CHANGE	0.055218
NAME_INCOME_TYPE_Working	0.057481
REGION_RATING_CLIENT	0.058899
REGION_RATING_CLIENT_W_CITY	0.060893
DAYS_EMPLOYED	0.074958
DAYS_BIRTH	0.078239
TARGET	1.000000

Name: TARGET, dtype: float64



Most Negative Correlations:

EXT_SOURCE_3	-0.178919
EXT_SOURCE_2	-0.160472
EXT_SOURCE_1	-0.155317
NAME_EDUCATION_TYPE_Higher education	-0.056593
CODE_GENDER_F	-0.054704
NAME_INCOME_TYPE_Pensioner	-0.046209
DAYS_EMPLOYED_ANOM	-0.045987
ORGANIZATION_TYPE_XNA	-0.045987
FLOORSMAX_AVG	-0.044003
FLOORSMAX_MEDI	-0.043768
FLOORSMAX_MODE	-0.043226
EMERGENCYSTATE_MODE_No	-0.042201
HOUSETYPE_MODE_block of flats	-0.040594
AMT_GOODS_PRICE	-0.039645
REGION_POPULATION_RELATIVE	-0.037227

Name: TARGET, dtype: float64

```
In [10]: # Find the correlation of the positive days since birth and target
app_train['DAYS_BIRTH'] = abs(app_train['DAYS_BIRTH'])
app_train['DAYS_BIRTH'].corr(app_train['TARGET'])
```

Out[10]: -0.07823930830982712

```
In [11]: # Age information into a separate dataframe
age_data = app_train[['TARGET', 'DAYS_BIRTH']]
age_data['YEARS_BIRTH'] = age_data['DAYS_BIRTH'] / 365

# Bin the age data
age_data['YEARS_BINNED'] = pd.cut(age_data['YEARS_BIRTH'], bins = np.linspace(20, 70,
age_data.head(10)
```

Out[11]:	TARGET	DAYS_BIRTH	YEARS_BIRTH	YEARS_BINNED
0	1	9461	25.920548	(25.0, 30.0]
1	0	16765	45.931507	(45.0, 50.0]
2	0	19046	52.180822	(50.0, 55.0]
3	0	19005	52.068493	(50.0, 55.0]
4	0	19932	54.608219	(50.0, 55.0]
5	0	16941	46.413699	(45.0, 50.0]
6	0	13778	37.747945	(35.0, 40.0]
7	0	18850	51.643836	(50.0, 55.0]
8	0	20099	55.065753	(55.0, 60.0]
9	0	14469	39.641096	(35.0, 40.0]

```
In [12]: # Group by the bin and calculate averages
age_groups = age_data.groupby('YEARS_BINNED').mean()
age_groups
```

```
Out [12]:
```

YEARS_BINNED	TARGET	DAYS_BIRTH	YEARS_BIRTH
(20.0, 25.0]	0.123036	8532.795625	23.377522
(25.0, 30.0]	0.111436	10155.219250	27.822518
(30.0, 35.0]	0.102814	11854.848377	32.479037
(35.0, 40.0]	0.089414	13707.908253	37.555913
(40.0, 45.0]	0.078491	15497.661233	42.459346
(45.0, 50.0]	0.074171	17323.900441	47.462741
(50.0, 55.0]	0.066968	19196.494791	52.593136
(55.0, 60.0]	0.055314	20984.262742	57.491131
(60.0, 65.0]	0.052737	22780.547460	62.412459
(65.0, 70.0]	0.037270	24292.614340	66.555108

```
In [13]: # Extract the EXT_SOURCE variables and show correlations
```

```
ext_data = app_train[['TARGET', 'EXT_SOURCE_1', 'EXT_SOURCE_2', 'EXT_SOURCE_3', 'DAYS_BIRTH']]
ext_data_corrs = ext_data.corr()
ext_data_corrs
```

```
Out [13]:
```

	TARGET	EXT_SOURCE_1	EXT_SOURCE_2	EXT_SOURCE_3	DAYS_BIRTH
TARGET	1.000000	-0.155317	-0.160472	-0.178919	-0.078239
EXT_SOURCE_1	-0.155317	1.000000	0.213982	0.186846	0.600610
EXT_SOURCE_2	-0.160472	0.213982	1.000000	0.109167	0.091996
EXT_SOURCE_3	-0.178919	0.186846	0.109167	1.000000	0.205478
DAYS_BIRTH	-0.078239	0.600610	0.091996	0.205478	1.000000

```
In [14]: # Make a new dataframe for polynomial features
```

```
poly_features = app_train[['EXT_SOURCE_1', 'EXT_SOURCE_2', 'EXT_SOURCE_3', 'DAYS_BIRTH']]
poly_features_test = app_test[['EXT_SOURCE_1', 'EXT_SOURCE_2', 'EXT_SOURCE_3', 'DAYS_BIRTH']]
```

```
# imputer for handling missing values
```

```
from sklearn.preprocessing import Imputer
imputer = Imputer(strategy = 'median')
```

```
poly_target = poly_features['TARGET']
```

```
poly_features = poly_features.drop(columns = ['TARGET'])
```

```
# Need to impute missing values
```

```
poly_features = imputer.fit_transform(poly_features)
poly_features_test = imputer.transform(poly_features_test)
```

```
from sklearn.preprocessing import PolynomialFeatures
```

```
# Create the polynomial object with specified degree
```

```
poly_transformer = PolynomialFeatures(degree = 3)
```

```
In [15]: # Train the polynomial features
```

```
poly_transformer.fit(poly_features)
```

```

# Transform the features
poly_features = poly_transformer.transform(poly_features)
poly_features_test = poly_transformer.transform(poly_features_test)
print('Polynomial Features shape: ', poly_features.shape)

```

Polynomial Features shape: (307511, 35)

```

In [16]: poly_transformer.get_feature_names(input_features = ['EXT_SOURCE_1', 'EXT_SOURCE_2',

```

```

Out[16]: ['1',
          'EXT_SOURCE_1',
          'EXT_SOURCE_2',
          'EXT_SOURCE_3',
          'DAYS_BIRTH',
          'EXT_SOURCE_1^2',
          'EXT_SOURCE_1 EXT_SOURCE_2',
          'EXT_SOURCE_1 EXT_SOURCE_3',
          'EXT_SOURCE_1 DAYS_BIRTH',
          'EXT_SOURCE_2^2',
          'EXT_SOURCE_2 EXT_SOURCE_3',
          'EXT_SOURCE_2 DAYS_BIRTH',
          'EXT_SOURCE_3^2',
          'EXT_SOURCE_3 DAYS_BIRTH',
          'DAYS_BIRTH^2']

```

```

In [17]: # Create a dataframe of the features

```

```

poly_features = pd.DataFrame(poly_features,
                             columns = poly_transformer.get_feature_names(['EXT_SOURCE_1', 'EXT_SOURCE_2', 'EXT_SOURCE_3', 'DAYS_BIRTH']))

```

```

# Add in the target

```

```

poly_features['TARGET'] = poly_target

```

```

# Find the correlations with the target

```

```

poly_corrs = poly_features.corr()['TARGET'].sort_values()

```

```

# Display most negative and most positive

```

```

print(poly_corrs.head(10))

```

```

print(poly_corrs.tail(5))

```

```

EXT_SOURCE_2 EXT_SOURCE_3          -0.193939
EXT_SOURCE_1 EXT_SOURCE_2 EXT_SOURCE_3  -0.189605
EXT_SOURCE_2 EXT_SOURCE_3 DAYS_BIRTH    -0.181283
EXT_SOURCE_2^2 EXT_SOURCE_3          -0.176428
EXT_SOURCE_2 EXT_SOURCE_3^2          -0.172282
EXT_SOURCE_1 EXT_SOURCE_2          -0.166625
EXT_SOURCE_1 EXT_SOURCE_3          -0.164065
EXT_SOURCE_2                     -0.160295

```

```

EXT_SOURCE_2 DAYS_BIRTH          -0.156873
EXT_SOURCE_1 EXT_SOURCE_2^2      -0.156867
Name: TARGET, dtype: float64
DAYS_BIRTH      -0.078239
DAYS_BIRTH^2    -0.076672
DAYS_BIRTH^3    -0.074273
TARGET          1.000000
1              NaN
Name: TARGET, dtype: float64

```

```

In [18]: # Put test features into dataframe
poly_features_test = pd.DataFrame(poly_features_test,
                                   columns = poly_transformer.get_feature_names(['EXT_SOURCE_1', 'EXT_SOURCE_2', 'EXT_SOURCE_3']))

# Merge polynomial features into training dataframe
poly_features['SK_ID_CURR'] = app_train['SK_ID_CURR']
app_train_poly = app_train.merge(poly_features, on = 'SK_ID_CURR', how = 'left')

# Merge polynomial features into testing dataframe
poly_features_test['SK_ID_CURR'] = app_test['SK_ID_CURR']
app_test_poly = app_test.merge(poly_features_test, on = 'SK_ID_CURR', how = 'left')

# Align the dataframes
app_train_poly, app_test_poly = app_train_poly.align(app_test_poly, join = 'inner', axis = 1)

# Print out the new shapes
print('Training data with polynomial features shape: ', app_train_poly.shape)
print('Testing data with polynomial features shape: ', app_test_poly.shape)

```

```

Training data with polynomial features shape:  (307511, 275)
Testing data with polynomial features shape:   (48744, 275)

```

```

In [19]: app_train_domain = app_train.copy()
app_test_domain = app_test.copy()

app_train_domain['CREDIT_INCOME_PERCENT'] = app_train_domain['AMT_CREDIT'] / app_train_domain['AMT_ANNUITY']
app_train_domain['ANNUITY_INCOME_PERCENT'] = app_train_domain['AMT_ANNUITY'] / app_train_domain['AMT_CREDIT']
app_train_domain['CREDIT_TERM'] = app_train_domain['AMT_CREDIT'] / app_train_domain['AMT_ANNUITY']
app_train_domain['DAYS_EMPLOYED_PERCENT'] = app_train_domain['DAYS_EMPLOYED'] / app_train_domain['DAYS_EMPLOYED']

app_test_domain['CREDIT_INCOME_PERCENT'] = app_test_domain['AMT_CREDIT'] / app_test_domain['AMT_ANNUITY']
app_test_domain['ANNUITY_INCOME_PERCENT'] = app_test_domain['AMT_ANNUITY'] / app_test_domain['AMT_CREDIT']
app_test_domain['CREDIT_TERM'] = app_test_domain['AMT_CREDIT'] / app_test_domain['AMT_ANNUITY']
app_test_domain['DAYS_EMPLOYED_PERCENT'] = app_test_domain['DAYS_EMPLOYED'] / app_test_domain['DAYS_EMPLOYED']

```

```

In [20]: from sklearn.ensemble import RandomForestClassifier

```

```

# Make the random forest classifier
random_forest = RandomForestClassifier(n_estimators = 100, random_state = 50, verbose=

```

In [13]: `from sklearn.preprocessing import MinMaxScaler, Imputer`

```

# Drop the target from the training data
if 'SK_ID_CURR' in app_train:
    train = app_train.drop(columns = ['SK_ID_CURR'])
else:
    train = app_train.copy()

```

```

# Feature names
features = list(train.columns)

```

```

# Copy of the testing data
test = app_test.copy()

```

```

# Median imputation of missing values
imputer = Imputer(strategy = 'median')

```

```

# Scale each feature to 0-1
scaler = MinMaxScaler(feature_range = (0, 1))

```

```

# Fit on the training data
imputer.fit(train)

```

```

# Transform both training and testing data
train = imputer.transform(train)
test = imputer.transform(app_test)

```

```

# Repeat with the scaler
scaler.fit(train)
train = scaler.transform(train)
test = scaler.transform(test)

```

```

print('Training data shape: ', train.shape)
print('Testing data shape: ', test.shape)

```

Training data shape: (356255, 46)

Testing data shape: (48744, 46)

In [49]: `# Train on the training data`  
`random_forest.fit(train, train_labels)`

```

# Extract feature importances
feature_importance_values = random_forest.feature_importances_
feature_importances = pd.DataFrame({'feature': features, 'importance': feature_importances_})

```

```

# Make predictions on the test data
predictions = random_forest.predict_proba(test)[: , 1]

[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed: 32.1s
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed: 1.4min finished
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed: 0.2s
[Parallel(n_jobs=8)]: Done 100 out of 100 | elapsed: 0.9s finished

```

```

In [50]: # Make a submission dataframe
submit = app_test[['SK_ID_CURR']]
submit['SK_ID_CURR'] = predictions

```

```

In [51]: submit.head(10)

```

```

Out[51]:   SK_ID_CURR  TARGET
0      100001      0.13
1      100005      0.21
2      100013      0.05
3      100028      0.14
4      100038      0.19
5      100042      0.15
6      100057      0.12
7      100065      0.14
8      100066      0.11
9      100067      0.18

```

```

In [26]: from sklearn.naive_bayes import GaussianNB

```

```

# Perform naive Bayes classification
nb = GaussianNB()
nb.fit(train.data, train_labels.data)

```

```

#feature_importance_values = nb.feature_importances_
#feature_importances = pd.DataFrame({'feature': features, 'importance': feature_impor

nb_res = nb.predict_proba(test.data)[: , 1]

```

```

In [27]: # Make a submission dataframe
submit1 = app_test[['SK_ID_CURR']]
submit1['TARGET'] = nb_res

```

```

In [28]: submit1.head(10)

```

```

Out[28]:   SK_ID_CURR      TARGET
0      100001  1.000000e+00
1      100005  1.000000e+00

```

```

2      100013  1.000000e+00
3      100028  1.935577e-08
4      100038  1.000000e+00
5      100042  9.401605e-01
6      100057  1.000000e+00
7      100065  1.000000e+00
8      100066  1.000000e+00
9      100067  1.000000e+00

```

```

In [52]: from sklearn.cross_validation import train_test_split
        from sklearn.tree import DecisionTreeClassifier
        from sklearn.metrics import accuracy_score
        from sklearn import tree

```

```

In [63]: clf_gini = DecisionTreeClassifier(random_state = 50, min_samples_leaf=1, min_samples_
        clf_gini.fit(train, train_labels)

```

```

Out[63]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
        max_features=None, max_leaf_nodes=None,
        min_impurity_decrease=0.0, min_impurity_split=None,
        min_samples_leaf=1, min_samples_split=2,
        min_weight_fraction_leaf=0.0, presort=False, random_state=50,
        splitter='best')

```

```

In [64]: feature_importance_values = clf_gini.feature_importances_
        feature_importances = pd.DataFrame({'feature': features, 'importance': feature_importa

        # Make predictions on the test data
        prediction1 = clf_gini.predict_proba(test)[: , 1]

```

```

In [65]: # Make a submission dataframe
        submit = app_test[['SK_ID_CURR']]
        submit['TARGET'] = prediction1

```

```

In [66]: submit.head(20)

```

```

Out[66]:      SK_ID_CURR  TARGET
0      100001      0.0
1      100005      1.0
2      100013      0.0
3      100028      0.0
4      100038      1.0
5      100042      1.0
6      100057      1.0
7      100065      0.0
8      100066      0.0
9      100067      0.0
10     100074      0.0
11     100090      0.0

```

12	100091	0.0
13	100092	0.0
14	100106	0.0
15	100107	1.0
16	100109	0.0
17	100117	0.0
18	100128	0.0
19	100141	0.0

```
In [ ]: submit.to_csv("first.csv", index=False)
```

```
In [33]: # Applying Linear Discriminant Analysis
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
lda = LinearDiscriminantAnalysis(n_components = 2)
lda.fit(train, train_labels)

X_test = lda.predict_proba(test)[: , 1]
```

```
-----
ValueError                                Traceback (most recent call last)
```

```
<ipython-input-33-762a387016f4> in <module>()
      2 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
      3 lda = LinearDiscriminantAnalysis(n_components = 2)
----> 4 lda.fit(train, train_labels)
      5
      6 X_test = lda.predict_proba(test)[: , 1]

~\Anaconda3\lib\site-packages\sklearn\discriminant_analysis.py in fit(self, X, y)
427         Target values.
428         """
--> 429         X, y = check_X_y(X, y, ensure_min_samples=2, estimator=self)
430         self.classes_ = unique_labels(y)
431
```

```
~\Anaconda3\lib\site-packages\sklearn\utils\validation.py in check_X_y(X, y, accept_sparse)
581         y = y.astype(np.float64)
582
--> 583         check_consistent_length(X, y)
584
585         return X, y
```

```
~\Anaconda3\lib\site-packages\sklearn\utils\validation.py in check_consistent_length(*arrays)
```



```

202     if len(uniques) > 1:
203         raise ValueError("Found input variables with inconsistent numbers of"
--> 204             " samples: %r" % [int(l) for l in lengths])
205
206

```

ValueError: Found input variables with inconsistent numbers of samples: [356255, 30751]

```

In [34]: submit = app_test[['SK_ID_CURR']]
        submit['TARGET'] = X_test

```

NameError Traceback (most recent call last)

```

<ipython-input-34-9a95cc9a64ea> in <module>()
    1 submit = app_test[['SK_ID_CURR']]
----> 2 submit['TARGET'] = X_test

```

NameError: name 'X\_test' is not defined

```

In [35]: submit.head()

```

```

Out[35]:   SK_ID_CURR
0      100001.0
1      100005.0
2      100013.0
3      100028.0
4      100038.0

```

```

In [23]: app_train_domain = app_train_domain.drop(columns = 'TARGET')

```

```

domain_features_names = list(app_train_domain.columns)

```

```

# Impute the domainnomial features
imputer = Imputer(strategy = 'median')

```

```

domain_features = imputer.fit_transform(app_train_domain)
domain_features_test = imputer.transform(app_test_domain)

```

```

# Scale the domainnomial features
scaler = MinMaxScaler(feature_range = (0, 1))

```

```

domain_features = scaler.fit_transform(domain_features)

```

```

domain_features_test = scaler.transform(domain_features_test)

#random_forest_domain = RandomForestClassifier(n_estimators = 100, random_state = 50,
XGB_params = {'num_round':200}
xgb_domain = XGBClassifier(max_depth = 6, learning_rate=0.2, estimator =100, **XGB_p

# Train on the training data
#random_forest_domain.fit(domain_features, train_labels)
xgb_domain.fit(domain_features, train_labels)

# Extract feature importances
#feature_importance_values_domain = random_forest_domain.feature_importances_
#feature_importances_domain = pd.DataFrame({'feature': domain_features_names, 'import

# Make predictions on the test data
#predictions = random_forest_domain.predict_proba(domain_features_test)[:, 1]
X_test = xgb_domain.predict_proba(domain_features_test)[:, 1]

```

-----

KeyError Traceback (most recent call last)

```

<ipython-input-23-d86b07a564f8> in <module>()
----> 1 app_train_domain = app_train_domain.drop(columns = 'TARGET')
      2
      3 domain_features_names = list(app_train_domain.columns)
      4
      5 # Impute the domainnomial features

~\Anaconda3\lib\site-packages\pandas\core\frame.py in drop(self, labels, axis, index,
3692                                     index=index, columns=columns,
3693                                     level=level, inplace=inplace,
-> 3694                                     errors=errors)
3695
3696     @rewrite_axis_style_signature('mapper', [('copy', True),

~\Anaconda3\lib\site-packages\pandas\core\generic.py in drop(self, labels, axis, index
3106         for axis, labels in axes.items():
3107             if labels is not None:
-> 3108                 obj = obj._drop_axis(labels, axis, level=level, errors=errors)
3109
3110             if inplace:

~\Anaconda3\lib\site-packages\pandas\core\generic.py in _drop_axis(self, labels, axis,

```

```

3138             new_axis = axis.drop(labels, level=level, errors=errors)
3139         else:
-> 3140             new_axis = axis.drop(labels, errors=errors)
3141             dropped = self.reindex(**{axis_name: new_axis})
3142             try:

~\Anaconda3\lib\site-packages\pandas\core\indexes\base.py in drop(self, labels, errors)
4385         if errors != 'ignore':
4386             raise KeyError(
-> 4387                 'labels %s not contained in axis' % labels[mask])
4388             indexer = indexer[~mask]
4389             return self.delete(indexer)

```

KeyError: "labels ['TARGET'] not contained in axis"

```

In [32]: submit = app_test[['SK_ID_CURR']]
         submit['TARGET'] = X_test

```

```

In [ ]: submit.head()

```

```

In [1]: # Applying XGB
        from xgboost import XGBClassifier
        XGB_params = {'num_round':200}
        xgb = XGBClassifier(max_depth = 6, learning_rate=0.2, estimator =100, **XGB_params)
        xgb.fit(train, train_labels)

        X_test = xgb.predict_proba(test)[:, 1]

```

-----

NameError Traceback (most recent call last)

```

<ipython-input-1-1d451edc2c0f> in <module>()
      3 XGB_params = {'num_round':200}
      4 xgb = XGBClassifier(max_depth = 6, learning_rate=0.2, estimator =100, **XGB_params)
----> 5 xgb.fit(train, train_labels)
      6
      7 X_test = xgb.predict_proba(test)[:, 1]

```

NameError: name 'train' is not defined

```

In [28]: submit = app_test[['SK_ID_CURR']]
         submit['TARGET'] = X_test

```

```
In [29]: submit.head()
```

```
Out[29]:
```

	SK_ID_CURR	TARGET
0	100001	0.028452
1	100005	0.091167
2	100013	0.015969
3	100028	0.012882
4	100038	0.069109